

Autoscaling How-To

This How-To expects that you have Terraform set-up successfully with a non-local state and kubectl connected with a Kubernetes Cluster.

In this example we begin with a CCE instance with one default node pool that has 2 nodes.

Create Autoscaling Node pool with Terraform

```
resource "opentelekomcloud_cce_node_pool_v3" "node_pool_autoscale" {
  cluster_id      = var.cluster_id
  name            = "${var.cluster_name}-node-pool-autoscale"
  flavor          = var.node_flavor
  initial_node_count = 1
  availability_zone = var.availability_zone
  key_pair        = var.key_pair_id
  os              = "EulerOS 2.5"

  scale_enable      = true
  min_node_count    = 1
  max_node_count    = 4
  scale_down_cooldown_time = 30 # minutes
  priority          = 1

  root_volume {
    size      = 100
    volumetype = "SATA"
  }
  data_volumes {
    size      = 100
    volumetype = "SATA"
  }
}
```

Apply this script with Terraform. Make sure that the availability zone, cluster_id and key_pair_id match the settings of your CCE. The node flavor can differ from the nodes in your default node pool and should be chosen with the specific workloads in mind.

After Terraform yields success, verify success in Web Console:

The screenshot shows the Open Telekom Cloud Web Console interface. The top navigation bar includes the Open Telekom Cloud logo, a language dropdown set to 'English', and a user profile icon. The left sidebar contains a menu with options like Dashboard, Workloads, Resource Management, Clusters, Nodes, Node Pools (highlighted), Network, Storage, and Namespaces. The main content area is titled 'Node Pools' and features a search bar and a 'Create Node Pool' button. Below this, a notification banner indicates that auto scaling is enabled. Two node pool cards are displayed: 'otc-custom...' with 1 node, 2 cores, and 8 GB memory, and 'DefaultPool' with 2 nodes, 4 cores, and 16 GB memory. Both pools are in a 'Normal' state.

Create Cluster Autoscaler Addon with Terraform

```
resource "opentelekomcloud_cce_addon_v3" "autoscaler" {
  template_name      = "autoscaler"
  template_version   = "1.17.2"
  cluster_id         = var.cluster_id

  values {
    basic = {
      "cceEndpoint": "https://cce.eu-de.otc.t-systems.com",
      "ecsEndpoint": "https://ecs.eu-de.otc.t-systems.com",
      "euleros_version": "2.2.5",
      "region": "eu-de",
      "swr_addr": "100.125.7.25:20202",
      "swr_user": "hwofficial"
    }
    custom = {
      "cluster_id": var.cluster_id,
      "tenant_id": var.project_id,
      # Max scaled up cores.
      "coresTotal": 32000,
      # Maximum number of empty nodes that can be deleted at the same time.
      "maxEmptyBulkDeleteFlag": 10,
      # Max scaled up nodes.
      "maxNodesTotal": 1000,
      # Max scaled up memory(GB).
      "memoryTotal": 128000,
      # The time after scale-up that the scale-down evaluation will resume.
      "scaleDownDelayAfterAdd": 10, # minutes
      # The time after node deletion that the scale-down evaluation will resume.
      "scaleDownDelayAfterDelete": 10, # minutes
      # The time after a scale-down failure that the scale-down evaluation will
      resume.
      "scaleDownDelayAfterFailure": 3, # minutes
      "scaleDownEnabled": true,
    }
  }
}
```

```
# The time of node is not used, default 10min.
"scaleDownUnneededTime": 10,
# The percent resource of node used to scale down.
"scaleDownUtilizationThreshold": 0.25,
"scaleUpCpuUtilizationThreshold": 0.8,
"scaleUpMemUtilizationThreshold": 0.8,
# Scale up when a pod cannot be scheduled because of missing resources
"scaleUpUnscheduledPodEnabled": true,
# Scale up when the utilization thresholds above are exceeded.
"scaleUpUtilizationEnabled": true,
# The timeout before autoscaler checks again the node that could not be
previously removed.
"unremovableNodeRecheckTimeout": 5 # minutes
    }
  }
}
```

A few things to check:

- The tenant_id to give to the autoscaler add-on must be the Project ID that you can find in the Web Console under "IAM" and "Projects"
- You can enable scale-up and scale-down independently
- If your loads are fluctuating frequently, you should be careful about not scale down too soon.

Again, we verify the success in the Web Console:

The screenshot shows the Open Telekom Cloud Web Console interface. The top navigation bar includes the logo, 'OPEN TELEKOM CLOUD', and a search bar. The left sidebar contains a menu with options like Dashboard, Workloads, Resource Management, Charts, Add-ons (highlighted), Auto Scaling, Permissions Management, Configuration Center, and Image Repository. The main content area is titled 'Add-ons' and shows a list of installed add-ons for the cluster 'otc-customer-succ...'. The add-ons are:

Add-on Name	Status	Version	Type	Actions
metrics-server	Running	1.0.3	Maintenance	Upgrade, Uninstall
everest	Running	1.1.11	Storage	Upgrade, Uninstall
coreDNS	Running	1.6.5	Service Discovery	Upgrade, Uninstall
autoscaler	Running	1.17.2	Maintenance	Upgrade, Uninstall

Testing Scaling up and down

We first test the scaling up by adding a test deployment:

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: autoscale-test-deployment
  labels:
    app: autoscale-test
spec:
  replicas: 1
  selector:
    matchLabels:
      app: autoscale-test
  template:
    metadata:
      labels:
        app: autoscale-test
    spec:
      containers:
        - name: hello-world
          image: nginx
          ports:
            - containerPort: 80
          resources:
            requests:
              memory: "64Mi"
              cpu: "250m"

```

We can scale the deployment and see how the cluster responds:

```
> kubectl scale deployment/autoscale-test-deployment --replicas=40
```

Since the 40 replicas utilize 10 CPUs, these do not fit on the nodes in the default node pool. Therefore the autoscaler will kick in and create an additional node.

```
> kubectl get pods
```

NAME	READY	STATUS	RESTARTS
AGE			
autoscale-test-deployment-6f9ff6448-4x248	0/2	Pending	0
13s			
autoscale-test-deployment-6f9ff6448-5kdcn	0/2	PodInitializing	0
14s			
autoscale-test-deployment-6f9ff6448-6pcmv	0/2	Pending	0
14s			
autoscale-test-deployment-6f9ff6448-8ftc8	1/2	Running	0
14s			
autoscale-test-deployment-6f9ff6448-9kxvt	0/2	Pending	0
14s			
autoscale-test-deployment-6f9ff6448-9scj5	0/2	Pending	0
13s			

autoscale-test-deployment-6f9ff6448-d7btf13s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-dsrvs14s	0/2	PodInitializing	0
autoscale-test-deployment-6f9ff6448-dxf5814s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-gdjvx14s	0/2	PodInitializing	0
autoscale-test-deployment-6f9ff6448-grws114s	0/2	PodInitializing	0
autoscale-test-deployment-6f9ff6448-gxbr913s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-h27z214s	0/2	Init:0/1	0
autoscale-test-deployment-6f9ff6448-h89vw13s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-hltfb13s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-hs5q813s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-m5zn914s	0/2	PodInitializing	0
autoscale-test-deployment-6f9ff6448-m6fxx14s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-mmtz214s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-mrpjt14s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-mzkrn26h	2/2	Running	0
autoscale-test-deployment-6f9ff6448-n6hrq14s	1/2	Running	0
autoscale-test-deployment-6f9ff6448-p2p9v14s	0/2	PodInitializing	0
autoscale-test-deployment-6f9ff6448-pt4vj13s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-q2ksm14s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-q7p7t14s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-qfbqq13s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-qs94913s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-qszsx14s	0/2	PodInitializing	0
autoscale-test-deployment-6f9ff6448-rm6c914s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-rnfzn14s	0/2	PodInitializing	0
autoscale-test-deployment-6f9ff6448-rsgh614s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-sgzhb13s	0/2	Pending	0

autoscale-test-deployment-6f9ff6448-v8qvm 13s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-w57gp 14s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-wfp5p 14s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-xh5sm 13s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-xrnrz 13s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-z72sp 13s	0/2	Pending	0
autoscale-test-deployment-6f9ff6448-zdgkp 14s	0/2	PodInitializing	0

In the Web Console, we see that the Scale Up event takes place:

The screenshot shows the Open Telekom Cloud Web Console interface. The left sidebar contains navigation options like Dashboard, Workloads, Resource Management, Charts, Add-ons, Auto Scaling, Permissions Management, Configuration Center, and Image Repository. The main content area displays the 'autoscaler' configuration for the cluster 'otc-customer-success-dev'. It shows the status as 'Running', version 1.17.2, and a recent 'Install complete' event. Below this, the 'Events' tab is selected, showing a table of events. The table has columns for Event Name, Kubernetes Event, First Occurred, and Last Occurred. A recent event is listed: 'Scale-up: setting group 137ef88d-a11e-11eb-8691-0255ac1017b0 size to 4'.

Event Name	Kubernetes Event	First Occurred	Last Occurred
Scale-up: setting group 137ef88d-a11e-11eb-8691-0255ac1017b0 size to 4		2021-04-21 09:10:02 GMT+02:00	2021-04-21 13:29:26 GMT+02:00





And then we also see the started nodes, 2 in the default node pool and 4 in the autoscale node pool:


```
> kubectl get nodes -L cce.cloud.com/cce-nodpool
NAME                                STATUS    ROLES    AGE    VERSION
CCE-NODEPOOL
192.168.13.187                      Ready    <none>    6m23s  v1.17.9-r0-CCE20.7.1.B003-17.36.3
otc-customer-success-dev-node-pool-autoscale
192.168.161.247                    Ready    <none>    4h15m  v1.17.9-r0-CCE20.7.1.B003-17.36.3
otc-customer-success-dev-node-pool-autoscale
192.168.182.115                    Ready    <none>    39d    v1.17.9-r0-CCE20.7.1.B003-17.36.3
192.168.186.181                    Ready    <none>    6m23s  v1.17.9-r0-CCE20.7.1.B003-17.36.3
otc-customer-success-dev-node-pool-autoscale
192.168.42.133                     Ready    <none>    39d    v1.17.9-r0-CCE20.7.1.B003-17.36.3
192.168.83.154                     Ready    <none>    6m17s  v1.17.9-r0-CCE20.7.1.B003-17.36.3
otc-customer-success-dev-node-pool-autoscale
```

Scaling down again...

```
> kubectl scale deployment/autoscale-test-deployment --replicas=1
```

yields removed nodes in the Web Console.

 **OPEN TELEKOM CLOUD** | Homepage | eu-de | English |  |  | 



Dashboard

Workloads

Resource Management

Charts

Add-ons


Auto Scaling **NEW**

Permissions Management

Configuration Center

Image Repository

autoscaler



Cluster Name

otc-customer-success-dev

Version

Ver. 1.17.2

Status

Running

Latest Event

Upgrade complete

Installed

2021-04-20 17:44:55 GMT+02:00

Updated

2021-04-22 11:00:59 GMT+02:00

Upgrade Add-on

Uninstall

Resources

Parameters


Events

Event data is retained for 1 hour and automatically cleared.


Start Date


—

End Date



Enter a Kubernetes event.





Even...	Occurre...	Event Name	Kubernetes Event	First Occurred	Last Occurred
Normal	1	Normal event	Scale-down: removing empty node 192.168.107.194	2021-04-22 12:00:49 GMT+02:00	2021-04-22 12:00:49 GMT+02:00